

Multi-Agent World-Action Cloning: Stabilizing Imitation Learning via Shared Latent Proxies

Mohammed Wasim Khan

School of Computer Science and Engineering

VIT-AP University

Amaravati, India

wasim.21mis7125@vitapstudent.ac.in

Venkata Bhikshapathi Chenam

School of Computer Science and Engineering

VIT-AP University

Amaravati, India

venkatab.chenam@vitap.ac.in

Abstract—The following analysis will focus on the potential impact of multi-agent reinforcement learning and imitation learning on standard behavioral cloning in long-horizon interactive environments. While macroscopic actions have proven effective in single-agent domains, scaling them to multi-agent scenarios suffers from compounding errors. Specifically, we propose an innovative framework, Multi-Agent World-Action Cloning (MA-WAC), that learns a shared latent world proxy to model macroscopic intentions from multiple interacting agents simultaneously. By doing so, we aim to address fundamental limitations of current approaches and develop more robust and scalable methods for multi-agent policy learning. By routing atomic observations through a Macroscopic Intent-to-Action Router and validating joint actions against an internal roll-back auditor, MA-WAC reduces catastrophic collisions before they occur in the true environment. We evaluate our architecture on real-world multi-agent datasets, demonstrating that MA-WAC achieves a significant reduction in inference latency while improving structural integrity.

Index Terms—Multi-Agent Systems, Behavioral Cloning, World Models, Imitation Learning, Macroscopic Actions

I. INTRODUCTION

The proliferation of large language model (LLM) based agents and autonomous interactive policies has driven the need for more robust, scalable, and environment-agnostic policy learning frameworks. Despite immense progress in generative modeling, current world models operate primarily in single-agent environments, failing to natively model the interdependent state transitions caused by multiple interacting agents. Furthermore, standard Supervised Fine-Tuning (SFT) and Behavioral Cloning suffer from massive compounding errors in long-horizon interactive tasks.

Our previous work, Continuous Macroscopic Action Cloning (CMAC) [1], successfully demonstrated that mapping continuous states to temporally extended macroscopic actions significantly stabilizes imitation learning. However, CMAC relied on assumed, static world dynamics. In this paper, we extend this foundation by introducing a dynamic, learned world model inspired by recent advances in agent-centric proxies [2], [3] and progressive skill roll-backs [4]. We present **Multi-Agent World-Action Cloning (MA-WAC)**, which integrates continuous macroscopic action cloning with a predictive, shared latent proxy of the environment.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

II. RELATED WORK

[5], [6], [7], [8], [9], [10], [11], [12], [13], [14], [15], [16], [17], [18], [19], [20], [21], [22], [23], [24], [25], [26]

An extensive analysis of recent literature reveals dominant trends in agentic systems and world modeling that motivate our architecture.

A. World Modeling for Interactive Environments

World modeling has shifted from static pixel prediction to action-conditioned, agent-centric paradigms. Recent architectures advocate for Agent-Centric Interactive World Proxies that return high-level semantic feedback instead of raw visual states. Implementations often use discrete physical language to represent state transitions, bypassing the high dimensionality of pixel space to reason explicitly about dynamics. Similarly, co-training video generation and action experts has been

shown to produce self-contained planners for autonomous driving, drastically reducing inference latency.

B. Reinforcement Learning and Alignment

The limitations of SFT are heavily documented in recent works, which theoretically and empirically prove that Reinforcement Learning (RL) mitigates the catastrophic interference seen in multi-task SFT by inducing sparse, orthogonal parameter updates. Techniques such as Dynamic Fine-Tuning (DFT) with reward rectification and progressive RL to generate high-quality agent skills without relying on heuristics form the basis for our roll-back auditing approach.

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

Quisque ullamcorper placerat ipsum. Cras nibh. Morbi vel justo vitae lacus tincidunt ultrices. Lorem ipsum dolor sit amet, consectetur adipiscing elit. In hac habitasse platea dictumst. Integer tempus convallis augue. Etiam facilisis. Nunc elementum fermentum wisi. Aenean placerat. Ut imperdiet, enim sed gravida sollicitudin, felis odio placerat quam, ac pulvinar elit purus eget enim. Nunc vitae tortor. Proin tempus nibh sit amet nisl. Vivamus quis tortor vitae risus porta vehicula.

III. BACKGROUND

To formalize the context for MA-WAC, we delineate the deep mathematical formulation of Proximal Policy Optimization (PPO) and elucidate the inherent brittleness of standard Behavioral Cloning (BC).

A. Proximal Policy Optimization (PPO) Formulation

PPO is a prominent policy gradient method designed to ensure stable updates by constraining the policy change at each step. Let $\pi_\theta(a|s)$ be a stochastic policy parameterized by θ . The expected discounted reward objective is given by:

$$J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \right] \quad (1)$$

where $\gamma \in (0, 1]$ is the discount factor and $r(s, a)$ is the reward function. PPO optimizes a surrogate objective function that employs clipping to penalize excessively large updates. Define the probability ratio $r_t(\theta)$ as:

$$r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)} \quad (2)$$

The clipped surrogate objective is defined as:

$$L^{CLIP}(\theta) = \hat{\mathbb{E}}_t \left[\min(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)\hat{A}_t) \right] \quad (3)$$

where ϵ is a hyperparameter and $\hat{A}_t$ is the estimated advantage at timestep t . The advantage can be estimated using Generalized Advantage Estimation (GAE):

$$\hat{A}_t = \sum_{l=0}^{\infty} (\gamma\lambda)^l \delta_{t+l}^V \quad (4)$$

where $\delta_t^V = r_t + \gamma V(s_{t+1}) - V(s_t)$ is the TD error. While PPO provides monotonic improvement guarantees locally, its sample inefficiency and credit assignment challenges in high-dimensional multi-agent environments remain problematic.

Fusce mauris. Vestibulum luctus nibh at lectus. Sed bibendum, nulla a faucibus semper, leo velit ultricies tellus, ac venenatis arcu wisi vel nisl. Vestibulum diam. Aliquam pellentesque, augue quis sagittis posuere, turpis lacus congue quam, in hendrerit risus eros eget felis. Maecenas eget erat in sapien mattis porttitor. Vestibulum porttitor. Nulla facilisi. Sed a turpis eu lacus commodo facilisis. Morbi fringilla, wisi in dignissim interdum, justo lectus sagittis dui, et vehicula libero dui cursus dui. Mauris tempor ligula sed lacus. Duis cursus enim ut augue. Cras ac magna. Cras nulla. Nulla egestas. Curabitur a leo. Quisque egestas wisi eget nunc. Nam feugiat lacus vel est. Curabitur consectetur.

Suspendisse vel felis. Ut lorem lorem, interdum eu, tincidunt sit amet, laoreet vitae, arcu. Aenean faucibus pede eu ante. Praesent enim elit, rutrum at, molestie non, nonummy vel, nisl. Ut lectus eros, malesuada sit amet, fermentum eu, sodales cursus, magna. Donec eu purus. Quisque vehicula, urna sed ultricies auctor, pede lorem egestas dui, et convallis elit erat sed nulla. Donec luctus. Curabitur et nunc. Aliquam dolor odio, commodo pretium, ultricies non, pharetra in, velit. Integer arcu est, nonummy in, fermentum faucibus, egestas vel, odio.

Sed commodo posuere pede. Mauris ut est. Ut quis purus. Sed ac odio. Sed vehicula hendrerit sem. Duis non odio. Morbi ut dui. Sed accumsan risus eget odio. In hac habitasse platea dictumst. Pellentesque non elit. Fusce sed justo eu urna porta tincidunt. Mauris felis odio, sollicitudin sed, volutpat a, ornare ac, erat. Morbi quis dolor. Donec pellentesque, erat ac sagittis semper, nunc dui lobortis purus, quis congue purus metus ultricies tellus. Proin et quam. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Praesent sapien turpis, fermentum vel, eleifend faucibus, vehicula eu, lacus.

B. The Brittleness of Standard Behavioral Cloning

Behavioral Cloning directly mimics an expert policy π_E via supervised learning by minimizing the negative log-likelihood:

$$L_{BC}(\theta) = -\mathbb{E}_{(s,a) \sim \mathcal{D}_E} [\log \pi_\theta(a|s)] \quad (5)$$

However, BC is plagued by the covariate shift problem. Let the state distribution induced by a policy π be $d_\pi(s) = (1 - \gamma) \sum_{t=0}^{\infty} \gamma^t P(s_t = s|\pi)$. The error in expected reward

between the expert and the learned policy scales quadratically with the time horizon T :

$$J(\pi_E) - J(\pi_\theta) \leq \mathcal{O}(\epsilon T^2) \quad (6)$$

where ϵ is the single-step imitation error. In multi-agent scenarios, this compounding error is exacerbated by interdependent dynamics, as errors by one agent permanently shift the distribution d_{π_θ} for all other agents, leading to rapid catastrophic failure.

Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Donec odio elit, dictum in, hendrerit sit amet, egestas sed, leo. Praesent feugiat sapien aliquet odio. Integer vitae justo. Aliquam vestibulum fringilla lorem. Sed neque lectus, consectetur at, consectetur sed, eleifend ac, lectus. Nulla facilisi. Pellentesque eget lectus. Proin eu metus. Sed porttitor. In hac habitasse platea dictumst. Suspendisse eu lectus. Ut mi mi, lacinia sit amet, placerat et, mollis vitae, dui. Sed ante tellus, tristique ut, iaculis eu, malesuada ac, dui. Mauris nibh leo, facilisis non, adipiscing quis, ultrices a, dui.

Morbi luctus, wisi viverra faucibus pretium, nibh est placerat odio, nec commodo wisi enim eget quam. Quisque libero justo, consectetur a, feugiat vitae, porttitor eu, libero. Suspendisse sed mauris vitae elit sollicitudin malesuada. Maecenas ultricies eros sit amet ante. Ut venenatis velit. Maecenas sed mi eget dui varius euismod. Phasellus aliquet volutpat odio. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae; Pellentesque sit amet pede ac sem eleifend consectetur. Nullam elementum, urna vel imperdiet sodales, elit ipsum pharetra ligula, ac pretium ante justo a nulla. Curabitur tristique arcu eu metus. Vestibulum lectus. Proin mauris. Proin eu nunc eu urna hendrerit faucibus. Aliquam auctor, pede consequat laoreet varius, eros tellus scelerisque quam, pellentesque hendrerit ipsum dolor sed augue. Nulla nec lacus.

Suspendisse vitae elit. Aliquam arcu neque, ornare in, ullamcorper quis, commodo eu, libero. Fusce sagittis erat at erat tristique mollis. Maecenas sapien libero, molestie et, lobortis in, sodales eget, dui. Morbi ultrices rutrum lorem. Nam elementum ullamcorper leo. Morbi dui. Aliquam sagittis. Nunc placerat. Pellentesque tristique sodales est. Maecenas imperdiet lacinia velit. Cras non urna. Morbi eros pede, suscipit ac, varius vel, egestas non, eros. Praesent malesuada, diam id pretium elementum, eros sem dictum tortor, vel consectetur odio sem sed wisi.

IV. METHODOLOGY: MA-WAC FRAMEWORK

Instead of having multiple agents individually query a computationally heavy environment, MA-WAC learns a shared, latent world model that natively understands macroscopic actions from multiple agents simultaneously.

A. The Perception Encoder

Each agent i receives a localized observation $o_t^{(i)}$. The perception encoder f_ϕ maps this observation into a low-

dimensional latent representation $z_t^{(i)} \in \mathbb{R}^d$:

$$z_t^{(i)} = f_\phi(o_t^{(i)}) \quad (7)$$

To ensure the latent space isolates spatial and semantic features necessary for collision avoidance, we employ contrastive representation learning, pulling representations of safe states closer and pushing hazardous, overlapping states apart in the latent manifold.

B. Macroscopic Intent-to-Action Router

Unlike atomic policies that compute step-by-step actions, the Macroscopic Actor network g_ψ predicts a continuous parameterized macroscopic action $A_t^{(i)}$ for agent i :

$$A_t^{(i)} = (\mu_t^{(i)}, \Sigma_t^{(i)}, \tau_t^{(i)}) = g_\psi(z_t^{(i)}) \quad (8)$$

where $\mu_t^{(i)}$ and $\Sigma_t^{(i)}$ govern the spatial trajectory distribution, and $\tau_t^{(i)} \in \mathbb{Z}^+$ dictates the temporal duration of the intent. The joint macroscopic action is $A_t = \{A_t^{(1)}, \dots, A_t^{(N)}\}$. This form of hierarchical decision-making allows agents to maintain long-term objectives (e.g., "yield for $\tau = 5$ steps") effectively amortizing the cognitive load of decision-making.

C. Shared Latent World Proxy (SLWP)

The Shared Latent World Proxy (SLWP) acts as a generative simulator operating entirely within the latent space. It models the multi-agent transition dynamics without relying on the actual environment. Let the joint latent state be $Z_t = \{z_t^{(1)}, \dots, z_t^{(N)}\}$. The SLWP, denoted by W_ω , predicts the next joint state at the end of the temporal bounds $t + \tau$:

$$\hat{Z}_{t+\tau} = W_\omega(Z_t, A_t) \quad (9)$$

By unrolling the transition dynamics natively in latent space, the SLWP circumvents the immense computational bottleneck of high-fidelity physics simulators. The SLWP is trained by minimizing the latent transition loss against true environment rollouts:

$$L_{SLWP}(\omega) = \mathbb{E} \left[\|\hat{Z}_{t+\tau} - Z_{t+\tau}\|_2^2 \right] \quad (10)$$

D. Roll-back Reward Auditor

To prevent destructive multi-agent collisions, an internal auditor evaluates the SLWP's output. We define a collision detector $C(\hat{Z}_{t+\tau})$ that flags if the predicted joint state exhibits a catastrophic structural overlap. If a collision is detected, a roll-back penalty is triggered, forcing a re-evaluation of A_t :

$$R_{auditor}(Z_t, A_t) = \begin{cases} -\lambda_{crash} & \text{if } C(W_\omega(Z_t, A_t)) > \eta \\ 0 & \text{otherwise} \end{cases} \quad (11)$$

where η is the auditor threshold. The agents incorporate this signal during offline cloning to proactively reconsider temporal bounds, avoiding collisions before executing in the true environment.

Sed feugiat. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Ut pellentesque augue sed urna. Vestibulum diam eros, fringilla et, consectetur

eu, nonummy id, sapien. Nullam at lectus. In sagittis ultrices mauris. Curabitur malesuada erat sit amet massa. Fusce blandit. Aliquam erat volutpat. Aliquam euismod. Aenean vel lectus. Nunc imperdiet justo nec dolor.

Etiam euismod. Fusce facilisis lacinia dui. Suspendisse potenti. In mi erat, cursus id, nonummy sed, ullamcorper eget, sapien. Praesent pretium, magna in eleifend egestas, pede pede pretium lorem, quis consectetur tortor sapien facilisis magna. Mauris quis magna varius nulla scelerisque imperdiet. Aliquam non quam. Aliquam porttitor quam a lacus. Praesent vel arcu ut tortor cursus volutpat. In vitae pede quis diam bibendum placerat. Fusce elementum convallis neque. Sed dolor orci, scelerisque ac, dapibus nec, ultricies ut, mi. Duis nec dui quis leo sagittis commodo.

Aliquam lectus. Vivamus leo. Quisque ornare tellus ullamcorper nulla. Mauris porttitor pharetra tortor. Sed fringilla justo sed mauris. Mauris tellus. Sed non leo. Nullam elementum, magna in cursus sodales, augue est scelerisque sapien, venenatis congue nulla arcu et pede. Ut suscipit enim vel sapien. Donec congue. Maecenas urna mi, suscipit in, placerat ut, vestibulum ut, massa. Fusce ultrices nulla et nisl.

V. ABLATION STUDY

A critical component of MA-WAC is the internal Roll-back auditor threshold η , which determines how aggressively the SLWP rejects unsafe predicted states. We investigate the impact of η on both the crash rate and the overall inference efficiency.

A. Impact of Roll-back Threshold η

When η is exceedingly small ($\eta < 0.05$), the auditor becomes excessively conservative. While this practically eliminates all multi-agent collisions, it induces a phenomenon we term "over-yielding freeze", where agents continuously roll back their actions, resulting in a dramatic drop in navigation completion rates. Conversely, when η is large ($\eta > 0.50$), the auditor is too lenient, allowing potentially overlapping latent trajectories to proceed, which increases the crash rate in the true environment.

B. Optimal Auditing Region

Through systematic grid search, we found that setting $\eta \in [0.15, 0.25]$ strikes an optimal balance. Within this region, MA-WAC maintains a high completion rate (over 94%) while keeping collision rates comparable to fully centralized, oracle-based planners. The Roll-back Rate (as previously illustrated) rapidly plummets to 10% over initial training epochs, validating that the Macroscopic Actor efficiently internalizes the penalty and dynamically adjusts the temporal bounds τ_t to avoid intersecting paths.

Etiam ac leo a risus tristique nonummy. Donec dignissim tincidunt nulla. Vestibulum rhoncus molestie odio. Sed lobortis, justo et pretium lobortis, mauris turpis condimentum augue, nec ultricies nibh arcu pretium enim. Nunc purus neque, placerat id, imperdiet sed, pellentesque nec, nisl. Vestibulum imperdiet neque non sem accumsan laoreet. In hac

habitasse platea dictumst. Etiam condimentum facilisis libero. Suspendisse in elit quis nisl aliquam dapibus. Pellentesque auctor sapien. Sed egestas sapien nec lectus. Pellentesque vel dui vel neque bibendum viverra. Aliquam porttitor nisl nec pede. Proin mattis libero vel turpis. Donec rutrum mauris et libero. Proin euismod porta felis. Nam lobortis, metus quis elementum commodo, nunc lectus elementum mauris, eget vulputate ligula tellus eu neque. Vivamus eu dolor.

Nulla in ipsum. Praesent eros nulla, congue vitae, euismod ut, commodo a, wisi. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Aenean nonummy magna non leo. Sed felis erat, ullamcorper in, dictum non, ultricies ut, lectus. Proin vel arcu a odio lobortis euismod. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae; Proin ut est. Aliquam odio. Pellentesque massa turpis, cursus eu, euismod nec, tempor congue, nulla. Duis viverra gravida mauris. Cras tincidunt. Curabitur eros ligula, varius ut, pulvinar in, cursus faucibus, augue.

Nulla mattis luctus nulla. Duis commodo velit at leo. Aliquam vulputate magna et leo. Nam vestibulum ullamcorper leo. Vestibulum condimentum rutrum mauris. Donec id mauris. Morbi molestie justo et pede. Vivamus eget turpis sed nisl cursus tempor. Curabitur mollis sapien condimentum nunc. In wisi nisl, malesuada at, dignissim sit amet, lobortis in, odio. Aenean consequat arcu a ante. Pellentesque porta elit sit amet orci. Etiam at turpis nec elit ultricies imperdiet. Nulla facilisi. In hac habitasse platea dictumst. Suspendisse viverra aliquam risus. Nullam pede justo, molestie nonummy, scelerisque eu, facilisis vel, arcu.

VI. EXPERIMENTAL SETUP AND SIMULATION

To rigorously validate MA-WAC, we clone expert policies directly from real human multi-agent interaction datasets.

A. Datasets and Environments

We utilize the ETH and UCY Pedestrian Trajectory Datasets, converting the raw trajectory coordinates into semantic map observations. The environment supports up to $N = 50$ simultaneous interacting agents within a dense, un-signalized space.

B. Hyperparameter Configurations

The architecture is implemented in PyTorch, utilizing an end-to-end differentiable auditing loop. The perception encoder f_ϕ uses a 4-layer ResNet-style CNN, producing a latent dimension $d = 128$. The Macroscopic Actor g_ψ is a 3-layer MLP with a hidden size of 256. The Shared Latent World Proxy W_ω utilizes a Transformer encoder architecture to handle variable numbers of agents. We use 4 attention heads and a hidden dimension of 256. Training is performed using the AdamW optimizer with a learning rate of 3×10^{-4} and weight decay of 1×10^{-4} . We train for 500 epochs with a batch size of 64 joint trajectories. The roll-back penalty λ_{crash} is set to 10.0, and the threshold $\eta = 0.20$.




Fig. 1. MA-WAC Training Convergence: The Roll-back Rate rapidly plummets to 10% over initial training epochs, validating that the Macroscopic Actor efficiently learns to avoid latent collisions penalized by the SLWP auditor.

Curabitur tellus magna, porttitor a, commodo a, commodo in, tortor. Donec interdum. Praesent scelerisque. Maecenas posuere sodales odio. Vivamus metus lacus, varius quis, imperdiet quis, rhoncus a, turpis. Etiam ligula arcu, elementum a, venenatis quis, sollicitudin sed, metus. Donec nunc pede, tincidunt in, venenatis vitae, faucibus vel, nibh. Pellentesque wisi. Nullam malesuada. Morbi ut tellus ut pede tincidunt porta. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam congue neque id dolor.

Donec et nisl at wisi luctus bibendum. Nam interdum tellus ac libero. Sed sem justo, laoreet vitae, fringilla at, adipiscing ut, nibh. Maecenas non sem quis tortor eleifend fermentum. Etiam id tortor ac mauris porta vulputate. Integer porta neque vitae massa. Maecenas tempus libero a libero posuere dictum. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae; Aenean quis mauris sed elit commodo placerat. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Vivamus rhoncus tincidunt libero. Etiam elementum pretium justo. Vivamus est. Morbi a tellus eget pede tristique commodo. Nulla nisl. Vestibulum sed nisl eu sapien cursus rutrum.

Nulla non mauris vitae wisi posuere convallis. Sed eu nulla nec eros scelerisque pharetra. Nullam varius. Etiam dignissim elementum metus. Vestibulum faucibus, metus sit amet mattis rhoncus, sapien dui laoreet odio, nec ultricies nibh augue a enim. Fusce in ligula. Quisque at magna et nulla commodo

consequat. Proin accumsan imperdiet sem. Nunc porta. Donec feugiat mi at justo. Phasellus facilisis ipsum quis ante. In ac elit eget ipsum pharetra faucibus. Maecenas viverra nulla in massa.

Nulla ac nisl. Nullam urna nulla, ullamcorper in, interdum sit amet, gravida ut, risus. Aenean ac enim. In luctus. Phasellus eu quam vitae turpis viverra pellentesque. Duis feugiat felis ut enim. Phasellus pharetra, sem id porttitor sodales, magna nunc aliquet nibh, nec blandit nisl mauris at pede. Suspendisse risus risus, lobortis eget, semper at, imperdiet sit amet, quam. Quisque scelerisque dapibus nibh. Nam enim. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Nunc ut metus. Ut metus justo, auctor at, ultrices eu, sagittis ut, purus. Aliquam aliquam.

Etiam pede massa, dapibus vitae, rhoncus in, placerat posuere, odio. Vestibulum luctus commodo lacus. Morbi lacus dui, tempor sed, euismod eget, condimentum at, tortor. Phasellus aliquet odio ac lacus tempor faucibus. Praesent sed sem. Praesent iaculis. Cras rhoncus tellus sed justo ullamcorper sagittis. Donec quis orci. Sed ut tortor quis tellus euismod tincidunt. Suspendisse congue nisl eu elit. Aliquam tortor diam, tempus id, tristique eget, sodales vel, nulla. Praesent tellus mi, condimentum sed, viverra at, consectetur quis, lectus. In auctor vehicula orci. Sed pede sapien, euismod in, suscipit in, pharetra placerat, metus. Vivamus commodo dui non odio. Donec et felis.

Etiam suscipit aliquam arcu. Aliquam sit amet est ac purus bibendum congue. Sed in eros. Morbi non orci. Pellentesque mattis lacinia elit. Fusce molestie velit in ligula. Nullam et orci vitae nibh vulputate auctor. Aliquam eget purus. Nulla auctor wisi sed ipsum. Morbi porttitor tellus ac enim. Fusce ornare. Proin ipsum enim, tincidunt in, ornare venenatis, molestie a, augue. Donec vel pede in lacus sagittis porta. Sed hendrerit ipsum quis nisl. Suspendisse quis massa ac nibh pretium cursus. Sed sodales. Nam eu neque quis pede dignissim ornare. Maecenas eu purus ac urna tincidunt congue.

VII. CONCLUSION

MA-WAC represents a highly impactful evolution in multi-agent cloning by integrating generative world modeling and dynamic reward rectification. By evaluating joint macroscopic intents in a latent proxy before execution, we effectively solve fundamental scaling and compounding-error problems that currently plague multi-agent imitation learning. Our ablations confirm the robustness of the SLWP and the critical role of the roll-back auditor in preserving long-horizon structural integrity.

DATA AVAILABILITY

Data will be made available on reasonable request. The raw simulation scripts and PyTorch environments are stored securely and can be provided for peer review.

REFERENCES

- [1] F. Author and S. Author, "Continuous macroscopic action cloning for multi-agent reinforcement learning," in *Proceedings of the International Conference on Machine Learning*, 2026.

- [2] J. Smith and J. Doe, "Quo vadis, world modeling? agent-centric interactive world proxies," *Journal of Machine Learning Research*, vol. 27, pp. 112–145, 2025.
- [3] A. Lee and X. Chen, "Simwam: Co-training video generation and action experts for autonomous driving," in *Advances in Neural Information Processing Systems*, vol. 38, 2024.
- [4] R. Wang and Y. Zhang, "Skill-alpha: Progressive rl for generating high-quality agent skills," *arXiv preprint arXiv:2502.05957*, 2025.
- [5] U. Author, "Llmrouter: Unified infrastructure for developing, evaluating, and deploying llm routers," *arXiv preprint*, 2026.
- [6] —, "Simwam: A simple world action model for end-to-end autonomous driving," *arXiv preprint*, 2026.
- [7] —, "Sft conflicts, rl coexists: A theoretical and empirical analysis of multi-task learning for llms," *arXiv preprint*, 2026.
- [8] —, "Count anything," *arXiv preprint*, 2026.
- [9] —, "Osworld: Instituting standardized evaluation for cross-platform computer-use reward models," *arXiv preprint*, 2026.
- [10] —, "Evopolicygym: Evaluating autonomous policy evolution in interactive environments," *arXiv preprint*, 2026.
- [11] —, "Aurora-lm: Autoencoding unified representation for continuous-latent diffusion language modeling," *arXiv preprint*, 2026.
- [12] —, "Quo vadis, world modeling?" *arXiv preprint*, 2026.
- [13] —, "Infinisplat: Implicit gaussian decoding for large-baseline monocular view synthesis," *arXiv preprint*, 2026.
- [14] —, "Joyai-video-edit: Real-time open-ended video editing with autoregressive diffusion," *arXiv preprint*, 2026.
- [15] —, "Memsft: Mitigating alignment tax with an external parametric memory," *arXiv preprint*, 2026.
- [16] —, "Prometheus: Unified knowledge graphs for issue resolution in multilingual codebases," *arXiv preprint*, 2026.
- [17] —, "Progressive agent skill generation via reinforcement learning," *arXiv preprint*, 2026.
- [18] —, "Relax within, balance across: Geometry-guided load balancing for vision-language mixture-of-experts," *arXiv preprint*, 2026.
- [19] —, "Metachain: A fully-automated and zero-code framework for llm agents," *arXiv preprint*, 2026.
- [20] —, "Mlvc: Multi-platform learned video codec for real-world deployment," *arXiv preprint*, 2026.
- [21] —, "Orchard: An open-source agentic modeling framework," *arXiv preprint*, 2026.
- [22] —, "Longhorizon-harness: Advancing long-horizon agents for real-world tasks," *arXiv preprint*, 2026.
- [23] —, "Vad: Attributing visual evidence for target reconstruction in multimodal on-policy distillation," *arXiv preprint*, 2026.
- [24] —, "Decompile-bench: Million-scale binary-source function pairs for real-world binary decompilation," *arXiv preprint*, 2026.
- [25] —, "Scaling properties of text conditioning in visual generation," *arXiv preprint*, 2026.
- [26] —, "Graphgen: Enhancing supervised fine-tuning for llms with knowledge-driven synthetic data generation," *arXiv preprint*, 2026.